

CAP-TTS-CGAN — Repo Skeleton

This canvas contains a ready-to-run skeleton for adapting **TTS-CGAN** to generate synthetic CAP data from the CAPSLPDB (PhysioNet).

Files included below (you can copy each into your project):

- `README.md` — quick start and overview
- `requirements.txt` — python deps
- `dataloader.py` — CAPDataset & preprocessing
- `models.py` — Generator and Discriminator (multichannel + optional conditioning)
- `train.py` — training script
- `generate.py` — sampling script
- `utils.py` — helper functions: annotation parsing, windowing, saving

README.md

```
# CAP-TTS-CGAN (skeleton)
```

```
This repository provides a skeleton for training a Time-series GAN (TTS-CGAN style) on the CAP Sleep Database (CAPSLPDB) to produce synthetic CAP EEG segments.
```

```
## Quick steps
```

```
1. Install dependencies:
```

```
```bash
pip install -r requirements.txt
```
```

```
2. Download CAPSLPDB from PhysioNet and place EDF + annotation files into a folder, e.g. `data/raw`.
```

```
3. Preprocess & create .npy dataset with provided dataloader utilities.
```

```
4. Train:
```

```
```bash
python train.py --data data/processed/cap_windows.npy --epochs 200 --batch_size 64
```
```

```
5. Generate samples:
```

```
```bash
python generate.py --model checkpoints/generator.pth --n 100 --out_dir generated
```
```

```
## Notes
```

- This skeleton is a starting point. Replace the annotation parsing and channel selection with channels you need (e.g., C3-A2, C4-A1).
- Conditioning is optional; there's an example of including label conditioning for CAP phase.

```
---

## requirements.txt
```

torch>=1.10 numpy pyedflib tqdm scipy mne

```
---

## dataloader.py

```python
import os
import numpy as np
import pyedflib
from torch.utils.data import Dataset

from utils import parse_cap_annotations, window_signal

class CAPDataset(Dataset):
 """Reads EDF files and corresponding annotation files, returns windows and
 optional labels.

 Expected input layout:
 data/raw/
 subject1.edf
 subject1_annotations.txt
 subject2.edf
 subject2_annotations.txt

 This dataset returns tensors shaped (channels, seq_len) per sample (PyTorch
 conv1d expects channels-first).
 """
 def __init__(self, raw_dir, channels=None, window_sec=30, fs=100,
use_labels=True, cache_npy=None):
 self.raw_dir = raw_dir
 self.window_sec = window_sec
 self.fs = fs
 self.use_labels = use_labels
 self.channels = channels # list of channel names or indices to select;
None = use all
 self.windows = [] # numpy arrays (channels, seq_len)
```

```

self.labels = [] # ints or None

if cache_npy and os.path.exists(cache_npy):
 print(f"Loading preprocessed data from {cache_npy}")
 data = np.load(cache_npy, allow_pickle=True)
 self.windows = data["windows"].tolist()
 self.labels = data["labels"].tolist()
else:
 self._build_dataset()
 if cache_npy:
 np.savez(cache_npy, windows=self.windows, labels=self.labels)

def _build_dataset(self):
 files = os.listdir(self.raw_dir)
 edf_files = [f for f in files if f.lower().endswith('.edf')]
 for edf in edf_files:
 edf_path = os.path.join(self.raw_dir, edf)
 base = os.path.splitext(edf)[0]
 ann_path = os.path.join(self.raw_dir, base + '_annotations.txt')

 try:
 f = pyedflib.EdfReader(edf_path)
 except Exception as e:
 print(f"Failed to read {edf_path}: {e}")
 continue

 n_signals = f.signals_in_file
 labels = f.getSignalLabels()
 sigs = np.array([f.readSignal(i) for i in range(n_signals)]) #
shape (n_channels, n_samples)
 f._close()
 del f

 # channel selection
 if self.channels is not None:
 # channels may be indices or names
 if all(isinstance(c, int) for c in self.channels):
 sigs = sigs[self.channels]
 else:
 idxs = [labels.index(c) for c in self.channels if c in
labels]
 sigs = sigs[idxs]

 # optionally parse annotations
 anns = None
 if self.use_labels and os.path.exists(ann_path):
 anns = parse_cap_annotations(ann_path)

```

```

 # windowing
 wlen = int(self.window_sec * self.fs)
 windows, win_labels = window_signal(sigs, anns, wlen, fs=self.fs)
 for win, lab in zip(windows, win_labels):
 # normalize per-window (z-score)
 win = (win - win.mean(axis=1, keepdims=True)) / (win.std(axis=1,
keepdims=True) + 1e-8)
 self.windows.append(win.astype(np.float32))
 self.labels.append(lab if lab is not None else -1)

 def __len__(self):
 return len(self.windows)

 def __getitem__(self, idx):
 x = self.windows[idx] # numpy (channels, seq_len)
 y = self.labels[idx]
 return x, y

```

## models.py

```

import torch
import torch.nn as nn

class ConditionalVector(nn.Module):
 """Simple embedding for label conditioning (if used).
 Returns a vector that can be concatenated to the latent vector or projected
 and added to feature maps.
 """
 def __init__(self, n_classes, latent_dim):
 super().__init__()
 self.embed = nn.Embedding(n_classes, latent_dim)
 def forward(self, labels):
 return self.embed(labels)

class Generator(nn.Module):
 def __init__(self, latent_dim=100, out_channels=3, seq_len=3000, cond=False,
n_classes=3):
 super().__init__()
 self.latent_dim = latent_dim
 self.cond = cond
 input_dim = latent_dim
 if cond:
 self.cond_vec = ConditionalVector(n_classes, latent_dim)

```

```

 # We'll map the latent vector into a set of feature maps and use
 ConvTranspose1d to upsample.
 self.project = nn.Sequential(
 nn.Linear(input_dim, 256 * (seq_len // 64)),
 nn.ReLU()
)
 self.deconv = nn.Sequential(
 nn.ConvTranspose1d(256, 128, kernel_size=4, stride=2, padding=1),
 nn.ReLU(),
 nn.ConvTranspose1d(128, 64, kernel_size=4, stride=2, padding=1),
 nn.ReLU(),
 nn.ConvTranspose1d(64, out_channels, kernel_size=4, stride=2,
padding=1),
 nn.Tanh()
)

 def forward(self, z, labels=None):
 if self.cond and labels is not None:
 cv = self.cond_vec(labels)
 z = z + cv
 x = self.project(z) # (batch, feat * time)
 batch = x.shape[0]
 feat_len = x.shape[1]
 # reshape to (batch, feat_maps, time)
 fmap = feat_len // (self.deconv[0].in_channels)
 # safer reshape: assume project set it to 256 * T
 x = x.view(batch, 256, -1)
 x = self.deconv(x)
 return x

class Discriminator(nn.Module):
 def __init__(self, in_channels=3, seq_len=3000, cond=False, n_classes=3):
 super().__init__()
 self.cond = cond
 if cond:
 self.label_proj = nn.Embedding(n_classes, in_channels)

 self.net = nn.Sequential(
 nn.Conv1d(in_channels, 64, kernel_size=7, stride=2, padding=3),
 nn.LeakyReLU(0.2),
 nn.Conv1d(64, 128, kernel_size=7, stride=2, padding=3),
 nn.LeakyReLU(0.2),
 nn.Conv1d(128, 256, kernel_size=7, stride=2, padding=3),
 nn.LeakyReLU(0.2),
 nn.AdaptiveAvgPool1d(1),
 nn.Flatten(),

```

```

 nn.Linear(256, 1)
)

 def forward(self, x, labels=None):
 # x shape: (batch, channels, seq_len)
 if self.cond and labels is not None:
 # simple conditioning by adding a per-channel bias from embedding
 emb = self.label_proj(labels) # shape (batch, in_channels)
 emb = emb.unsqueeze(-1) # (batch, in_channels, 1)
 x = x + emb
 out = self.net(x)
 return out

```

## utils.py

```

import numpy as np

def parse_cap_annotations(ann_path):
 """Parse CAP annotation file (stub).

 The exact format of annotations in CAPSLPDB varies. This function should be
 adapted to the real annotation layout.
 Expected return: list of tuples (onset_sample, duration_samples, label_int)
 """
 anns = []
 with open(ann_path, 'r') as f:
 for line in f:
 line = line.strip()
 if not line: continue
 # Example simple format: onset_sec,duration_sec,label
 parts = line.split(',')
 if len(parts) >= 3:
 onset = float(parts[0])
 dur = float(parts[1])
 lab = parts[2]
 # map lab to int (example)
 lab_int = 1 if 'A' in lab else 0
 anns.append((int(onset), int(dur), lab_int))
 return anns

def window_signal(sigs, anns, wlen, fs=100):
 """Slice multichannel signals into windows.

```

```

- sigs: np.array (n_channels, n_samples)
- anns: list of (onset_sec, dur_sec, label_int) OR None
- wlen: window length in samples

Returns (windows, labels)
 windows: list of np arrays (n_channels, wlen)
 labels: list of ints (e.g., majority label in window) or -1
"""
n_channels, n_samples = sigs.shape
step = wlen # non-overlapping by default; change to wlen//2 if overlap
desired
windows = []
labels = []
for start in range(0, n_samples - wlen + 1, step):
 win = sigs[:, start:start + wlen]
 label = -1
 if anns:
 # determine label by majority overlap
 start_sec = start / fs
 end_sec = (start + wlen) / fs
 counts = {}
 for a_on, a_dur, a_lab in anns:
 a_end = a_on + a_dur
 # check overlap in seconds
 if (a_on < end_sec) and (a_end > start_sec):
 counts[a_lab] = counts.get(a_lab, 0) + 1
 if counts:
 # choose label with max overlaps
 label = max(counts.items(), key=lambda x: x[1])[0]
 windows.append(win)
 labels.append(label)
return windows, labels

```

## train.py

```

import argparse
import torch
import torch.nn as nn
from torch.utils.data import DataLoader
import numpy as np
from tqdm import tqdm

from dataloader import CAPDataset
from models import Generator, Discriminator

```

```

def train(args):
 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

 # load dataset
 ds = CAPDataset(args.data_dir, window_sec=args.window_sec, fs=args.fs,
cache_npy=args.cache)
 # convert to DataLoader with collate to convert to torch tensors
 loader = DataLoader(ds, batch_size=args.batch_size, shuffle=True,
drop_last=True)

 # infer channels and seq_len from first sample
 sample_x, _ = ds[0]
 in_channels, seq_len = sample_x.shape

 G = Generator(latent_dim=args.latent_dim, out_channels=in_channels,
seq_len=seq_len).to(device)
 D = Discriminator(in_channels=in_channels, seq_len=seq_len).to(device)

 g_opt = torch.optim.Adam(G.parameters(), lr=args.lr, betas=(0.5, 0.999))
 d_opt = torch.optim.Adam(D.parameters(), lr=args.lr, betas=(0.5, 0.999))

 criterion = nn.BCEWithLogitsLoss()

 for epoch in range(args.epochs):
 pbar = tqdm(loader)
 for real, _ in pbar:
 real = torch.tensor(real).to(device) # numpy -> tensor
 # ensure shape (batch, channels, seq_len)
 if real.ndim == 3:
 pass
 else:
 real = real.unsqueeze(1)

 batch_size = real.size(0)

 # Train D
 z = torch.randn(batch_size, args.latent_dim).to(device)
 z = z
 fake = G(z)
 # make sure fake shape matches
 if fake.dim() == 2:
 fake = fake.unsqueeze(1)

 d_real = D(real)
 d_fake = D(fake.detach())
 real_labels = torch.ones_like(d_real)

```



```

 fake_labels = torch.zeros_like(d_fake)

 d_loss = criterion(d_real, real_labels) + criterion(d_fake,
fake_labels)
 d_opt.zero_grad(); d_loss.backward(); d_opt.step()

 # Train G
 d_fake_for_g = D(fake)
 g_loss = criterion(d_fake_for_g, real_labels)
 g_opt.zero_grad(); g_loss.backward(); g_opt.step()

 pbar.set_description(f"E{epoch} D:{d_loss.item():.4f} G:
{g_loss.item():.4f}")

 # checkpoint
 torch.save(G.state_dict(), f"checkpoints/generator_epoch{epoch}.pth")
 torch.save(D.state_dict(), f"checkpoints/
discriminator_epoch{epoch}.pth")

if __name__ == '__main__':
 parser = argparse.ArgumentParser()
 parser.add_argument('--data_dir', type=str, default='data/raw')
 parser.add_argument('--cache', type=str, default='data/processed/
cap_windows.npz')
 parser.add_argument('--window_sec', type=int, default=30)
 parser.add_argument('--fs', type=int, default=100)
 parser.add_argument('--epochs', type=int, default=50)
 parser.add_argument('--batch_size', type=int, default=32)
 parser.add_argument('--latent_dim', type=int, default=128)
 parser.add_argument('--lr', type=float, default=2e-4)
 args = parser.parse_args()
 train(args)

```

## generate.py

```

import argparse
import os
import torch
import numpy as np
from models import Generator

def generate(args):

```

```

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
load generator
G = Generator(latent_dim=args.latent_dim, out_channels=args.channels,
seq_len=args.seq_len).to(device)
G.load_state_dict(torch.load(args.model, map_location=device))
G.eval()

os.makedirs(args.out_dir, exist_ok=True)
n = args.n
batch = 64
generated = []
for i in range(0, n, batch):
 b = min(batch, n - i)
 z = torch.randn(b, args.latent_dim).to(device)
 with torch.no_grad():
 fake = G(z)
 fake = fake.cpu().numpy() # (batch, channels, seq_len)
 for j in range(fake.shape[0]):
 out = fake[j]
 np.save(os.path.join(args.out_dir, f"sample_{i+j}.npy"), out)
 generated.append(out)
print(f"Saved {len(generated)} samples to {args.out_dir}")

if __name__ == '__main__':
 parser = argparse.ArgumentParser()
 parser.add_argument('--model', type=str, required=True)
 parser.add_argument('--n', type=int, default=100)
 parser.add_argument('--out_dir', type=str, default='generated')
 parser.add_argument('--latent_dim', type=int, default=128)
 parser.add_argument('--channels', type=int, default=3)
 parser.add_argument('--seq_len', type=int, default=3000)
 args = parser.parse_args()
 generate(args)

```

## Final notes

- This skeleton aims to be a practical starting point. Replace the annotation parsing logic in `utils.parse_cap_annotations` with the real format used by CAPSLPDB.
- Tune model depth, kernel sizes, strides, and training hyperparameters for your dataset and sequence length.

<!-- End of canvas -->